

# SQL on Kafka

Usman Masood  
@usmanm



# What is PipelineDB?

- Open-source relational database
- Run SQL queries continuously on streaming data
- Integrates stream computation and persistent storage into a single system
- Fork of PostgreSQL 9.5
- Leverage existing clients and tools that support vanilla PostgreSQL

# Continuous Query Abstractions

- Continuous Views
- Continuous Transforms
- Continuous Triggers

# Continuous Views

- High throughput, real-time MATERIALIZED VIEW
- Data is aggregated before written to disk
- Raw granular stream data is discarded
- Database size is independent of the amount of data ingested over time
- Support for probabilistic structures that allow for efficient approximate computations over streams
- Support for sliding window queries

# Continuous Transforms

- Doesn't store the output of query, instead a trigger function is called for each output row
- Only non-aggregate queries are supported
- Normalize incoming data
- Push data to external systems

# Continuous Triggers

- Triggers on continuous views
- Any user-defined function can be used
- Can be called any time a row is inserted or updated in a continuous view
- Real-time anomaly detection and notifications

# pipeline\_kafka

- PostgreSQL extension written in C
- Uses `librdkafka`
- Consume messages from Kafka topics into streams or relations
- Durably snapshots offsets to state so messages are only consumed once
- Push `bytea` values or JSON serialized tuples into Kafka topics
- Already used in production by many of our users!

# pipeline\_kafka Broker API

- `pipeline_kafka.add_broker (host text)`
- `pipeline_kafka.remove_broker (host text)`



# pipeline\_kafka Consumer API

- `pipeline_kafka.consume_begin (`  
    `topic text,`  
    `relation text,`  
    `format := 'text',`  
    `delimiter := E'\t',`  
    `quote := NULL,`  
    `escape := NULL,`  
    `batchsize := 1000,`  
    `parallelism := 1,`  
    `start_offset := NULL`  
    `)`
- `pipeline_kafka.consume_end (topic text, relation text)`

# pipeline\_kafka Consumer API

- Each (topic, relation) gets parallelism number of processes
- Partitions are uniformly assigned to all processes
- Uses `rd_kafka_consume_batch` to consume messages in batches
- Uses PostgreSQL's `COPY` command to copy messages into relations or streams
- Update offsets after copying a batch

# pipeline\_kafka Producer API

- `pipeline_kafka.produce_message (`  
    `topic text,`  
    `message bytea,`  
    `partition := NULL,`  
    `key := NULL`  
    `)`
- `pipeline_kafka.emit_tuple ( topic, partition, key )`

# Putting It All Together

Team

id	name
1	Warriors
2	Thunder
3	Cavaliers
...	...

Player

id	team	name
1	1	Stephen Curry
2	1	Klay Thompson
3	2	Kevin Durant
...	...	...

# Putting It All Together

ppg\_topic

```
{ "id": 1, "points": 32 }
```

```
{ "id": 37, "points": 17 }
```

```
{ "id": 9, "points": 2 }
```

```
{ "id": 33, "points": 27 }
```

```
{ "id": 91, "points": 13 }
```

```
...
```

# Putting It All Together

pts\_season\_stats

id	name	team	avg	points	num_games
1	Stephen Curry	Warriors	30.06	2375	79
2	Klay Thompson	Warriors	22.13	1771	80
3	Kevin Durant	Thunder	28.18	2029	72
...	...	...	...	...	...

**Extra Credit:** Produce alert message into `mvp_topic` whenever a player crosses 1000 points scored in the season.

# Extract & Transform

ppg\_topic


```
{ "id": 1, "points": 32 }
```

```
{ "id": 37, "points": 17 }
```

```
{ "id": 9, "points": 2 }
```

```
{ "id": 33, "points": 27 }
```


...



ppg\_stream

id	name	team	points
1	Stephen Curry	Warriors	32
37	Chris Bosh	Heat	17
9	Lou Williams	Lakers	2
33	Tim Duncan	Spurs	27
...	...	...	...

# Transform & Load





# Slam Dunk With PipelineDB

Let's do it in 6 lines of SQL!

# Thanks!

- [www.pipelinedb.com](http://www.pipelinedb.com)
- [docs.pipelinedb.com](http://docs.pipelinedb.com)
- [github.com/pipelinedb](https://github.com/pipelinedb)
- @pipelinedb